

QRMumps.jl

QRMumps.jl is a Julia package designed to solve large, sparse linear systems of equations ($Ax=b$)

Technology

It acts as a wrapper for the QR_MUMPS library, which uses a multifrontal QR factorization method. This method is known for its numerical stability, especially with ill-conditioned matrices.

Audience

The software is targeted at computational scientists, engineers, and researchers using Julia for tasks that involve sparse matrices, such as numerical optimization.

Repository: github.com/JuliaQR/QRMumps.jl

Age: Started in 2017

Community: 6 lifetime contributors, with discussions on GitHub Issues

Activity Level: Mature and stable with a low commit frequency, which indicates a complete feature set rather than neglect.

Code Example

Setup

```
using SparseArrays, LinearAlgebra
```

```
# Small 4x4 sparse matrix
```

```
A = sparse([
    4.0  1.0  0.0  0.0;
    1.0  3.0  1.0  0.0;
    0.0  1.0  3.0  1.0;
    0.0  0.0  1.0  2.0
])
```

```
b = [15.0, 10.0, 10.0, 10.0]
```

```
display(A)
```

```
display(b)
```

```
4x4 SparseMatrixCSC{Float64, Int64} with 10 stored entries:
```

```
 4.0  1.0  .  .
 1.0  3.0  1.0  .
 .  1.0  3.0  1.0
 .  .  1.0  2.0
```

```
4-element Vector{Float64}:
```

```
15.0
10.0
10.0
10.0
```

Solve

```
using QRMumps, LinearAlgebra

qrm_init()

spmat = qrm_spmat_init(A)

spfct = qrm_spfct_init(spmat)

# Analyze sparsity, factorize, and solve
qrm_analyse!(spmat, spfct)
qrm_factorize!(spmat, spfct)
x = qrm_solve(spfct, b)

cond_number = cond(Matrix(A))

println("Solution x = ", x)
println("Residual norm = ", norm(A * x - b))
println("Condition Number = ", cond_number)

Solution x = [-3.001433187162955, 2.215685929711668, 5.988795543912927, -10.90406868485484]
Residual norm = 35.80278984429186
Condition Number = 4.2072442645144745
```

Question: How does QR factorization compare to the default LU for ill-conditioned systems?

Julia's default sparse solver uses LU factorization which is typically faster than QR.

However, QR factorization is known to be more numerically stable. For matrices that have a high condition number, this stability could be crucial for accuracy.

Proposed Experiment: Quantify the performance and accuracy trade-off between QRMumps.jl and Julia's default sparse solver.

Experimental Steps

1. **Generate Matrices:** Create a set of sparse matrices with increasing condition numbers. This systematically increases the difficulty of the problem.
2. **Solve:** For each matrix, solve the linear system $Ax=b$ using both methods
3. **Measure:** Record two key metrics for each solve: How long did it take, and how accurate is the solution?